

CS476 - PW4 Report

Julien de Castelnaud, Magnus Meyer

November 11, 2024

Exercises

1. `__packed` tells the compiler to ignore internal alignment requirements and the alignment of the struct itself.
2.
 - The ordering of the fields within the struct is fixed in all cases: 4 byte `id` followed directly by `data`, however long it is, plus padding at the end depending on `PARAM_DATALEN`
 - `id` is always 4 bytes, while `data` is `PARAM_DATALEN` bytes long.
 - If `__packed` is not present, padding bytes are `4 - PARAM_DATALEN` to satisfy requirement for `item_t` to be aligned to 4 byte boundary. Otherwise, it is 0.
3. If `__packed` is present, the `data` is immediately followed by the `id` of the next `item_t` in the array. Otherwise, there is the aforementioned number of bytes of padding separating the two elements.
4. The number of cache misses saturate at 256 when `PARAM_DATALEN > 24`. Since `PARAM_COUNT=256`, that means that there is a miss for every loop iteration. So every iteration accesses a different cache line, and since each iteration accesses the `i`th struct in the array, the struct must be at least as large as the cache line. The first size in which we saturate is when `PARAM_DATALEN = 25` or when the unpacked size is 32 bytes. So, the cache line size is 32 bytes.
5. `__packed` increases the generated code size, because the OpenRISC 32-bit load instruction doesn't handle unaligned addresses. Extra code needs to be added to assemble the number from individual bytes/halfwords.

Task 1

7. The accesses to `node->id` and `node->next` cause cache misses. They are always on different lines due to the size of `data`. The struct is 64 bytes in total, which is 2 cache lines, and `id` is at the start, while `next` is at the end.

TODO does the string literal cause a cache miss probably yes

8. Move `next` and `prev` to the start of the struct. This yields 25 dcache misses down from 43.

Task 2

9. The data cache misses come from accessing the `id` field of the struct at `items[i]`. Each iteration accesses a different line because the struct is 36 bytes long, while the cache line is 32 bytes.
10. ??? cant change `items_init` so idk

Store a pointer to the data instead of copying it. This reduces the size of the struct down to 8 bytes which should cut the number of misses roughly by a factor of 4.

Task 3

11. Row-major and column-major refers to the ordering of a 2D array in memory. In row-major, the row given by the first index is contiguous in memory. In column-major the column given by the second index is contiguous. C uses row-major.

12. Primarily the accesses to `out_vector` and `matrix`. Because `matrix` is row-major each inner iteration loads a different cache line (size of row is 256 elements or 1024 bytes). `out_vector` may be similarly evicted frequently because the entire 1KiB of vector is stored to on every outer loop iteration.
13. Reorder the loops so that `i` is on the outside and `j` inside.

Task 4

14. It does not work because the writes to the LEDS are being cached. Normally, the I/O should be in an uncached region of memory but this is not the case here.
15.
 - Disable the caches entirely by commenting out the body of `init_dcache()`
 - Disable write-back policy in `init_dcache()` and use write-through instead (remove `CACHE_WRITE_BACK` flag):

```
dcache_write_cfg (CACHE_FOUR_WAY | CACHE_SIZE_4K | CACHE_REPLACE_LRU);
```

This works because write-through means that the access is cached, but it is propagated to the system memory anyway.